


# Reinforcement learning

Bibliographic search by Matthieu Vilain

Tutor : Pierre Andry

May 2017


# What is reinforcement learning ?



Way of programming agents by reward and punishment without needing to specify how the task is to be achieved


- TD Learning - R.Sutton - 1988
- Q Learning - C.Watkins - 1989-1992

## Exemple




# How does it work ?

## The agent and his environment



The environment must be divided into several states

## Exemple

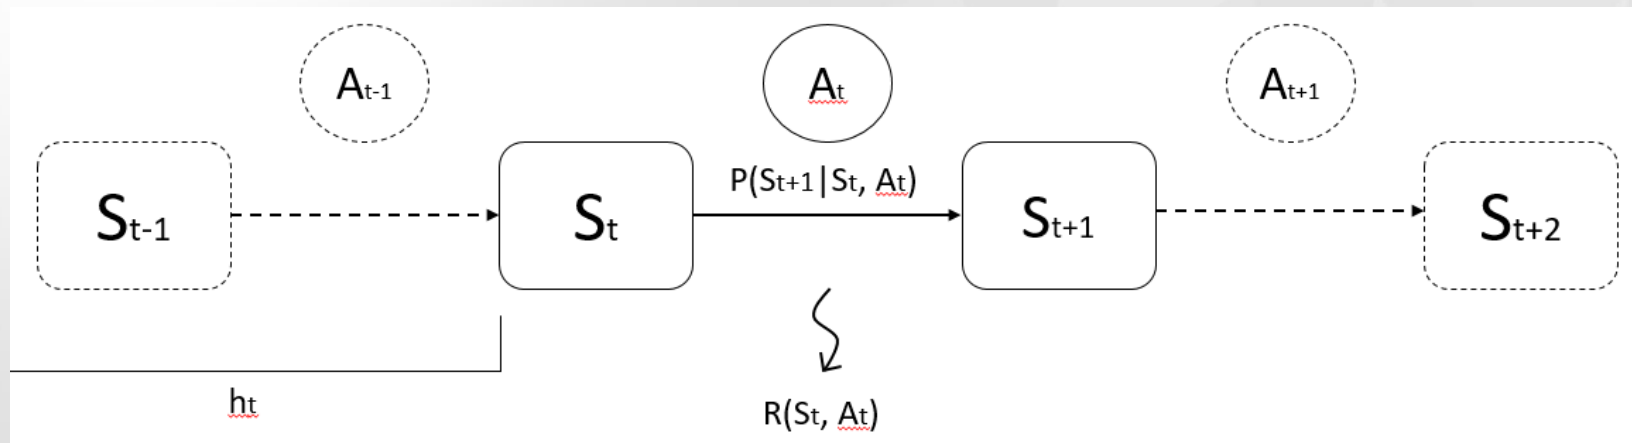


### Informations :

- List of possible actions
- Reward :
  - Positive
  - Negative
  - Null
- Value

## How does agent choose his action ?

### Markov Decision Process



- $S$  : states set
- $A$  : actions set
- $T$  : time

- $P()$  : probability of transition between states
- $R()$  : reward function for transition
- $h$  : historic of actions

# How does agent choose his action ?

## The policy

Definition :

Procedure to be followed by the agent to choose at a moment the action to be executed (noted  $\pi$ )

GOAL → found the optimal policy ( $\pi^*$ )

## Bellman equation

Estimation of the future reward if the agent do a action

$$V^\pi(s) = \sum_{a \in \mathcal{A}(s)} \pi(s, a) \sum_{s' \in \mathcal{S}} \mathcal{P}(s, a, s') [\mathcal{R}(s, a, s') + \gamma V^\pi(s')]$$

$\mathcal{P}(s, a, s')$  : probability to go to the state  $s'$  with the action  $a$  if I am in  $s$

[1] Sutton, R. Planning by incremental dynamic programming.

In *Eight International Workshop on Machine Learning*, pages 353-357. Morgan Kaufmann. 1991

[2] L.P. Kaelbling, M. L. Littman, A.W. Moore, « Reinforcement Learning : A Survey », *Journal of Artificial Intelligence Research* 4, 1996

[3] S. Russell P. Norving, Book « Artificial Intelligence : A Modern Approach » 2010

# Reinforcement learning algorithm : Q-Learning

The value of the estimation is not on the state  $S_t$  but on the action to go from  $S_t$  to  $S_{t+1}$

## From Bellman to Q-Learning

$$V^\pi(s) = \sum_{a \in \mathcal{A}(s)} \pi(s, a) \sum_{s' \in \mathcal{S}} P(s, a, s') [\mathcal{R}(s, a, s') + \gamma V^\pi(s')]$$



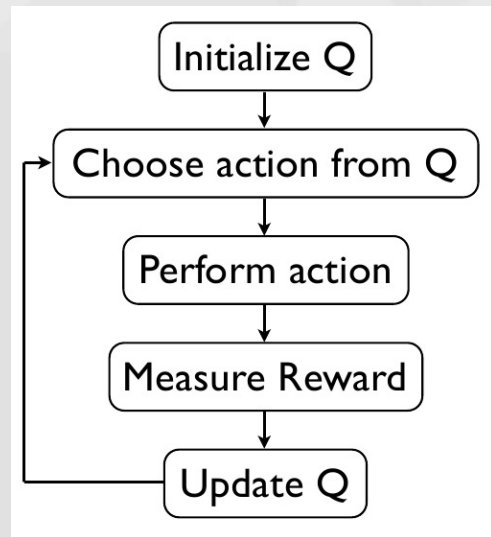
$$Q^*(s, a) = \sum_{s'} P(s, a, s') [\mathcal{R}(s, a, s') + \gamma \max_{a'} Q^*(s', a')]$$



Actualization function

$$Q(s_t, a_t) = Q(s_t, a_t) + \alpha (R_t + \gamma \max_a Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))$$


## Algorithm



# Q-Learning, algorithm explication

[0,0]	[0,1]	[0,2]	[0,3]	[0,4]	[0,5]	[0,6]	[0,7]	[0,8]	[0,9]
[1,0]	[1,1]	[1,2]	[1,3]	[1,4]	[1,5]	[1,6]	[1,7]	[1,8]	[1,9]
[2,0]	[2,1]	[2,2]	[2,3]	[2,4]	[2,5]	[2,6]	[2,7]	[2,8]	[2,9]
[3,0]	[3,1]	[3,2]	[3,3]	[3,4]	[3,5]	[3,6]	[3,7]	[3,8]	[3,9]
[4,0]	[4,1]	[4,2]	[4,3]	[4,4]	[4,5]	[4,6]	[4,7]	[4,8]	[4,9]
[5,0]	[5,1]	[5,2]	[5,3]	[5,4]	[5,5]	[5,6]	[5,7]	[5,8]	[5,9]
[6,0]	[6,1]	[6,2]	[6,3]	[6,4]	[6,5]	[6,6]	[6,7]	[6,8]	[6,9]
[7,0]	[7,1]	[7,2]	[7,3]	[7,4]	[7,5]	[7,6]	[7,7]	[7,8]	[7,9]
[8,0]	[8,1]	[8,2]	[8,3]	[8,4]	[8,5]	[8,6]	[8,7]	[8,8]	[8,9]
[9,0]	[9,1]	[9,2]	[9,3]	[9,4]	[9,5]	[9,6]	[9,7]	[9,8]	[9,9]

$$Q(s_t, a_t) = Q(s_t, a_t) + \alpha (R_t + \gamma \max_a Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))$$



1920 moves


19 backup

2mins (16moves/sec)

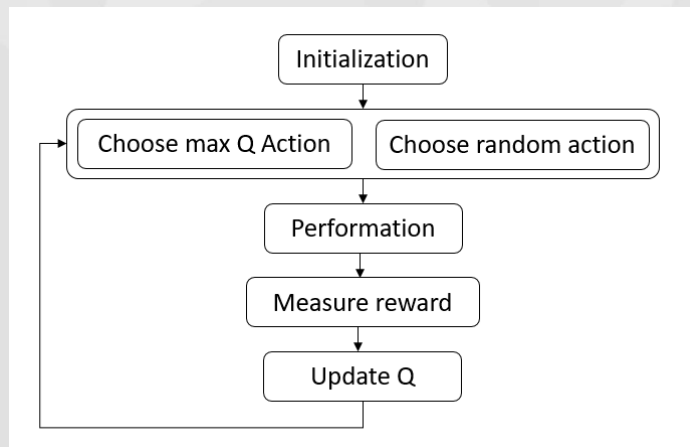
 $\alpha = 0,5$  $\gamma = 0,5$

# Issue exploration vs exploitation

## Issue

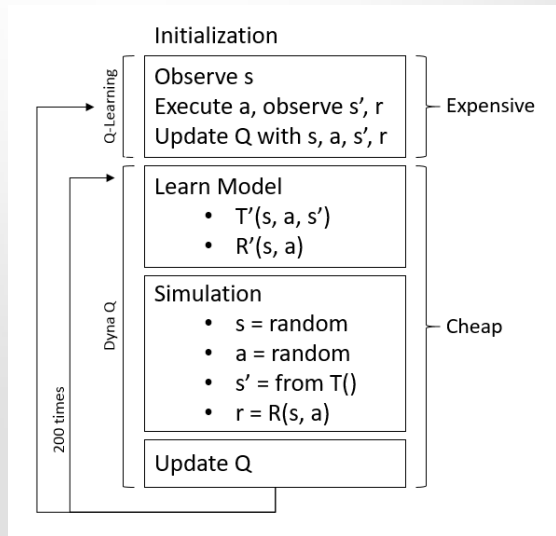


## Algorithm






# Algorithm



# Test




	Steps before convergence	Backups before convergence
Q-learning	531,000	531,000
Dyna	62,000	3,055,000
prioritized sweeping	28,000	1,010,000

- [1] L.P. Kaelbling, M. L. Littman, A.W. Moore, « Reinforcement Learning : A Survey », Journal of Artificial Intelligence Research 4, 1996  
 [2] Watkins, C.J.C.H. "Q-Learning", Machine Learning, 8(3), 279-292 , 1992  
 [3] H. Larochelle, YouTube video « Intelligence Artificielle [13.7] : Apprentissage par renforcement – Q-learning »  
 [4] P.Norving, S.Thrun, Udacity cours, « Introduction to Artificial Inteligence » 2010

In this project we will  
simulate the evolution of a  
population in a urban  
environment

# Reinforcement learning in robotics



To many states  $\Rightarrow$  exploration time  
Delimitate states in real world





Fig. 2. Five link biped robot

## Function approximator



- Make a generalization
- Use only the visible states

# Reinforcement learning

## Issues, discussion

**Lemma :** let  $n$  denote the number of actions applicable at state  $s'$ . If all  $n$  actions share target Q-Value, i.e,  
 $\exists q : \forall \hat{a} : q = Q^{target}(s', \hat{a})$ , then the average overestimation  $E[Z_s]$  is  $\gamma c$  with  $c = \varepsilon \frac{n-1}{n+1}$

**Table 1:** Upper bound on the error  $\varepsilon$  of the function approximator, according to Theorem 2. These bounds are significant. For example, if episodes of length  $L = 60$  with  $n = 5$  actions shall be learned,  $\varepsilon$  must be smaller than .00943 (bold number).

	$L=10$	$L=20$	$L=30$	$L=40$	$L=50$	$L=60$	$L=70$	$L=80$	$L=90$	$L=100$	$L=1000$
$n = 2$	.12991	.05966	.03872	.02866	.02275	.01886	.01611	.01405	.01247	.01120	.00110
$n = 3$	.08660	.03977	.02581	.01911	.01517	.01257	.01074	.00937	.00831	.00746	.00073
$n = 4$	.07217	.03314	.02151	.01592	.01264	.01048	.00895	.00781	.00692	.00622	.00061
$n = 5$	.06495	.02983	.01936	.01433	.01137	<b>.00943</b>	.00805	.00702	.00623	.00560	.00055
$n = 6$	.06062	.02784	.01807	.01337	.01061	.00880	.00751	.00656	.00581	.00522	.00051
$n = 8$	.05567	.02557	.01659	.01228	.00975	.00808	.00690	.00602	.00534	.00480	.00047
$n = 10$	.05292	.02430	.01577	.01167	.00927	.00768	.00656	.00572	.00508	.00456	.00045
$n = 20$	.04786	.02198	.01426	.01056	.00838	.00695	.00593	.00517	.00459	.00412	.00040
$n = \infty$	.04330	.01988	.01290	.00955	.00758	.00628	.00537	.00468	.00415	.00373	.00036

[1] J.Morimoto, G.Cheng, C.G.Atkeson, G.Zeglin, « A simple reinforcement learning algorithm for biped walking »

International conference on robotics & automation, New Orleans, 2004

[2] S.Thrun, A.Schwartz « Issues in using function approximation for reinforcement learning »

Proceedings of the Fourth Connectionist Models Summer School, Lawrence Erlbaum Publisher, Hillsdale, NJ, Dec. 1993

[3] L.P. Kaelbling, M. L. Littman, A.W. Moore, « Reinforcement Learning : A Survey », Journal of Artificial Intelligence Research 4, 1996

[4] S. Russell P. Norving, Book « Artificial Intelligence : A Modern Approach » 2010